

▼ Exploration of house data from Arizona state

```
# Import the packages
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import folium
```

▼ Data input

```
# Read the data
data = pd.read_csv('data_project.csv')
```

```
# Display the few rows of the data
data.head(5)
```

| | Unnamed: 0 | MLS | sold_price | zipcode | longitude | latitude | lot_acres | taxe |
|---|------------|----------|------------|---------|-------------|-----------|-----------|---------|
| 0 | 0 | 21530491 | 5300000.0 | 85637 | -110.378200 | 31.356362 | 2154.00 | 5272.0 |
| 1 | 1 | 21529082 | 4200000.0 | 85646 | -111.045371 | 31.594213 | 1707.00 | 10422.3 |
| 2 | 3 | 21919321 | 4500000.0 | 85646 | -111.035925 | 31.645878 | 636.67 | 8418.5 |
| 3 | 4 | 21306357 | 3411450.0 | 85750 | -110.813768 | 32.285162 | 3.21 | 15393.0 |
| 4 | 5 | 21528016 | 3250000.0 | 85718 | -110.910593 | 32.339090 | 1.67 | 27802.8 |

```
# Check the dimensions of the data
data.shape
```

```
(4903, 19)
```

```
data['sold_price'].min(), data['sold_price'].max()
```

```
(300000.0, 5300000.0)
```

```
# show the concise summary about data and count missing values in the dataframe
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4903 entries, 0 to 4902
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            4903 non-null   int64
1   MLS                                    4903 non-null   int64
2   sold_price                            4903 non-null   float64
3   zipcode                              4903 non-null   int64
4   longitude                             4903 non-null   float64
5   latitude                             4903 non-null   float64
6   lot_acres                             4903 non-null   float64
7   taxes                                 4903 non-null   float64
8   year_built                            4903 non-null   int64
9   bedrooms                              4903 non-null   int64
10  bathrooms                             4903 non-null   float64
11  sqrt_ft                               4903 non-null   float64
12  garage                                4903 non-null   float64
13  kitchen_features                      4903 non-null   object
14  fireplaces                            4903 non-null   int64
15  floor_covering                        4903 non-null   object
16  HOA                                    4903 non-null   object
17  sold_price_per_sqrt_ft                4903 non-null   float64
18  sold_price_per_sqrt_ft_range          4902 non-null   object
dtypes: float64(9), int64(6), object(4)
memory usage: 727.9+ KB
```

```
# Descriptive statistical measures
data.describe()
```

| | Unnamed: 0 | MLS | sold_price | zipcode | longitude | latitude |
|--------------|-------------|--------------|--------------|--------------|-------------|-------------|
| count | 4903.000000 | 4.903000e+03 | 4.903000e+03 | 4903.000000 | 4903.000000 | 4903.000000 |
| mean | 2495.650010 | 2.149141e+07 | 7.756850e+05 | 85724.291046 | -110.910560 | 32.315921 |
| std | 1443.203645 | 1.336633e+06 | 3.157269e+05 | 36.678685 | 0.120478 | 0.166568 |
| min | 0.000000 | 4.110917e+06 | 3.000000e+05 | 85118.000000 | -112.520168 | 31.356362 |
| 25% | 1245.500000 | 2.141039e+07 | 5.850000e+05 | 85718.000000 | -110.978280 | 32.280525 |
| 50% | 2494.000000 | 2.161781e+07 | 6.780000e+05 | 85737.000000 | -110.922309 | 32.319152 |
| 75% | 3745.500000 | 2.180551e+07 | 8.390000e+05 | 85750.000000 | -110.858404 | 32.396256 |
| max | 4998.000000 | 2.192856e+07 | 5.300000e+06 | 86323.000000 | -109.454637 | 34.927884 |

```
# check the duplicate values
data.duplicated().sum()
```

```
0
```

▼ Dealing with Missing values

```
# check the missing value in each column
data.isna().sum()
```

```
Unnamed: 0      0
MLS             0
sold_price      0
zipcode         0
longitude       0
latitude        0
lot_acres       0
taxes           0
year_built      0
bedrooms        0
bathrooms       0
sqrt_ft         0
garage          0
kitchen_features 0
fireplaces      0
floor_covering  0
HOA             0
sold_price_per_sqrt_ft 0
sold_price_per_sqrt_ft_range 1
dtype: int64
```

```
data1 = data.dropna()
data1.isna().sum()
```

```
Unnamed: 0      0
MLS             0
sold_price      0
zipcode         0
longitude       0
latitude        0
lot_acres       0
taxes           0
year_built      0
bedrooms        0
bathrooms       0
sqrt_ft         0
garage          0
kitchen_features 0
fireplaces      0
floor_covering  0
HOA             0
sold_price_per_sqrt_ft 0
sold_price_per_sqrt_ft_range 0
dtype: int64
```

```
data1["sold_price_per_sqrt_ft_range"] = data1["sold_price_per_sqrt_ft_range"].astype('category')
data1["sold_price_per_sqrt_ft_range"] = data1["sold_price_per_sqrt_ft_range"].cat.codes
data1.sold_price_per_sqrt_ft_range = data1.sold_price_per_sqrt_ft_range.astype(int)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:

```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html>
 """Entry point for launching an IPython kernel.
 /usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
 Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html>
 /usr/local/lib/python3.7/dist-packages/pandas/core/generic.py:5170: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
 Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html>
 self[name] = value

```
data2 = data1[['longitude', 'latitude', 'sold_price_per_sqrt_ft_range']]
data2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4902 entries, 0 to 4902
Data columns (total 3 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   longitude                            4902 non-null   float64
 1   latitude                             4902 non-null   float64
 2   sold_price_per_sqrt_ft_range         4902 non-null   int64
dtypes: float64(2), int64(1)
memory usage: 153.2 KB
```

```
"""
from sklearn.preprocessing import OrdinalEncoder
ord_enc = OrdinalEncoder()
data2["sold_price_per_sqrt_ft_range"] = ord_enc.fit_transform(data2[["sold_price_per_sqrt_ft_range"]])
data2.sold_price_per_sqrt_ft_range = data2.sold_price_per_sqrt_ft_range.astype(int)
"""
```

```
'\nfrom sklearn.preprocessing import OrdinalEncoder\nord_enc = OrdinalEncoder()\ndata2["sold_price_per_sqrt_ft_range"] = ord_enc.fit_transform(data2[["sold_price_per_sqrt_ft_range"]])\ndata2.sold_price_per_sqrt_ft_range = data2.sold_price_per_sqrt_ft_range.astype(int)
```

```
# Shuffle your dataset
shuffle_df = data2.sample(frac=1)
```

```
# Define a size for your train set
train_size = int(0.8 * len(data2))
```

```
# Split your dataset
train_set = shuffle_df[:train_size]
test_set = shuffle_df[train_size:]
```

```
X_train = train_set.drop(['sold_price_per_sqrt_ft_range'], axis=1)
```

```
X_train = X_train.to_numpy()
```

```
y_train=train_set.iloc[:,-1]
y_train = y_train.to_numpy()
y_train.shape
```

```
(3921,)
```

```
X_test = test_set.drop(['sold_price_per_sqrt_ft_range'], axis=1)
X_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 981 entries, 1837 to 953
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   longitude    981 non-null    float64
1   latitude     981 non-null    float64
dtypes: float64(2)
memory usage: 23.0 KB
```

```
X_test = X_test.to_numpy()
X_test.shape
X_test
```

```
array([[ -110.847125,  32.262861],
       [ -110.833476,  32.31409 ],
       [ -110.929826,  32.336238],
       ...,
       [ -111.02565 ,  32.434199],
       [ -111.039738,  32.274257],
       [ -110.702877,  32.271099]])
```

```
y_test=test_set.iloc[:,-1]
y_test = y_test.to_numpy()
y_test.shape
```

```
(981,)
```

```
class KNNClassifier():
    def fit(self, X, y):
        self.X=X
        self.y=y

    def predict(self, X,K, epsilon=10e-15):
        N=len(X)
```

```

y_hat = np.zeros(N)

for i in range(N):
    dist2 = np.sum((self.X-X[i])**2,axis=1)
    idxt = np.argsort(dist2)[:K]
    gamma_k = 1/(np.sqrt(dist2[idxt]+epsilon))
    y_hat[i] = np.bincount(self.y[idxt], weights=gamma_k).argmax()

return y_hat

knn = KNNClassifier()

knn.fit(X_train,y_train)

y_hat1 = knn.predict(X_train, 8)

y_hat = knn.predict(X_test, 8)

def accuracy(y, y_hat):
    return np.mean(y==y_hat)

accuracy(y_train, y_hat1)

0.9987248150981892

accuracy(y_test, y_hat)

0.6992864424057085

```

▼ Predicting the sold price

```

data3 = data1[['bedrooms', 'bathrooms','fireplaces','sold_price_per_sqrt_ft_range', 'sold_
data3.sold_price = data3.sold_price.astype(int)
data3

```



/usr/local/lib/python3.7/dist-packages/pandas/core/generic.py:5170: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/u>

self[name] = value

| | bedrooms | bathrooms | fireplaces | sold_price_per_sqrt_ft_range | sold_price |
|------|----------|-----------|------------|------------------------------|------------|
| 0 | 13 | 10.0 | 6 | 2 | 5300000 |
| 1 | 2 | 2.0 | 5 | 2 | 4200000 |
| 2 | 7 | 5.0 | 4 | 2 | 4500000 |
| 3 | 4 | 6.0 | 5 | 2 | 3411450 |
| 4 | 3 | 4.0 | 5 | 2 | 3250000 |
| 4898 | 3 | 2.0 | 1 | 0 | 5350000 |

Shuffle your dataset

```
shuffle_df1 = data3.sample(frac=1)
```

Define a size for your train set

```
train_size1 = int(0.8 * len(data3))
```

Split your dataset

```
train_set1 = shuffle_df1[:train_size1]
```

```
test_set1 = shuffle_df1[train_size1:]
```

```
X_train1 = train_set1.drop(['sold_price'], axis=1)
```

```
X_train1 = X_train1.to_numpy()
```

```
y_train1=train_set1.iloc[:,-1]
```

```
y_train1 = y_train1.to_numpy()
```

```
X_test1 = test_set1.drop(['sold_price'], axis=1)
```

```
X_test1 = X_test1.to_numpy()
```

```
y_test1=test_set1.iloc[:,-1]
```

```
y_test1 = y_test1.to_numpy()
```

```
def OLS(Y,Y_hat):
```

```
    N=Y.shape[0]
```

```
    return (1/(2*N)*np.sum((Y-Y_hat)**2))
```

```
class OurLinearRegression():
```

```
    def fit(self, X, y, eta=1e-3, epochs = 1e3, show_curve =False):
```

```
        epochs = int(epochs)
```

```
        N, D = X.shape
```

```
        Y = y
```

```

self.W = np.random.randn(D)
J = np.zeros(epochs)

for epoch in range(epochs):
    Y_hat = self.predict(X)
    J[epoch] = OLS(Y,Y_hat)
    self.W -= eta*(1/N)*(X.T@(Y_hat-Y))

```

```

if show_curve:
    plt.figure()
    plt.plot(J)
    plt.xlabel("epochs")
    plt.ylabel("$\text{match}")
    plt.title("Training ")

```


```

def predict(self, X):
    return X@self.W

```

```
myLireg = OurLinearRegression()
```

```
myLireg.fit(X_train1,y_train1, epochs = 3e4,eta = 1e-5, show_curve = True)
```



```

def R2(Y, Y_hat):
    return (1-(np.sum((Y-Y_hat)**2)/np.sum((Y-np.mean(Y))**2)))

```

```
y_hat_train = myLireg.predict(X_train1)
```

```
R2(y_train1,y_hat_train)
```

```
0.02515760208349438
```

```
y_hat_test = myLireg.predict(X_test1)
```

```
R2(y_test1,y_hat_test)
```



```
0.1699112273766934
```

```
X_t = np.array([[5,4, 5,5]])
```

```
myLireg.predict(X_t)
```

```
array([1080695.46738759])
```

```
d = data3.loc[data3['sold_price_per_sqrt_ft_range'] == 4]
d
```

| | bedrooms | bathrooms | fireplaces | sold_price_per_sqrt_ft_range | sold_price |
|-----------|----------|-----------|------------|------------------------------|------------|
| 13 | 3 | 3.0 | 3 | 4 | 3000000 |
| 14 | 3 | 3.0 | 3 | 4 | 2600000 |
| 17 | 3 | 3.0 | 3 | 4 | 2600000 |
| 19 | 3 | 3.0 | 3 | 4 | 2600000 |
| 25 | 3 | 3.0 | 3 | 4 | 2600000 |
| 28 | 2 | 2.0 | 1 | 4 | 2500000 |
| 47 | 2 | 3.0 | 1 | 4 | 1800000 |

```
data2['sold_price_per_sqrt_ft_range'].value_counts()
```

```
0    2516
1    2298
2     75
4     7
3     6
Name: sold_price_per_sqrt_ft_range, dtype: int64
```

